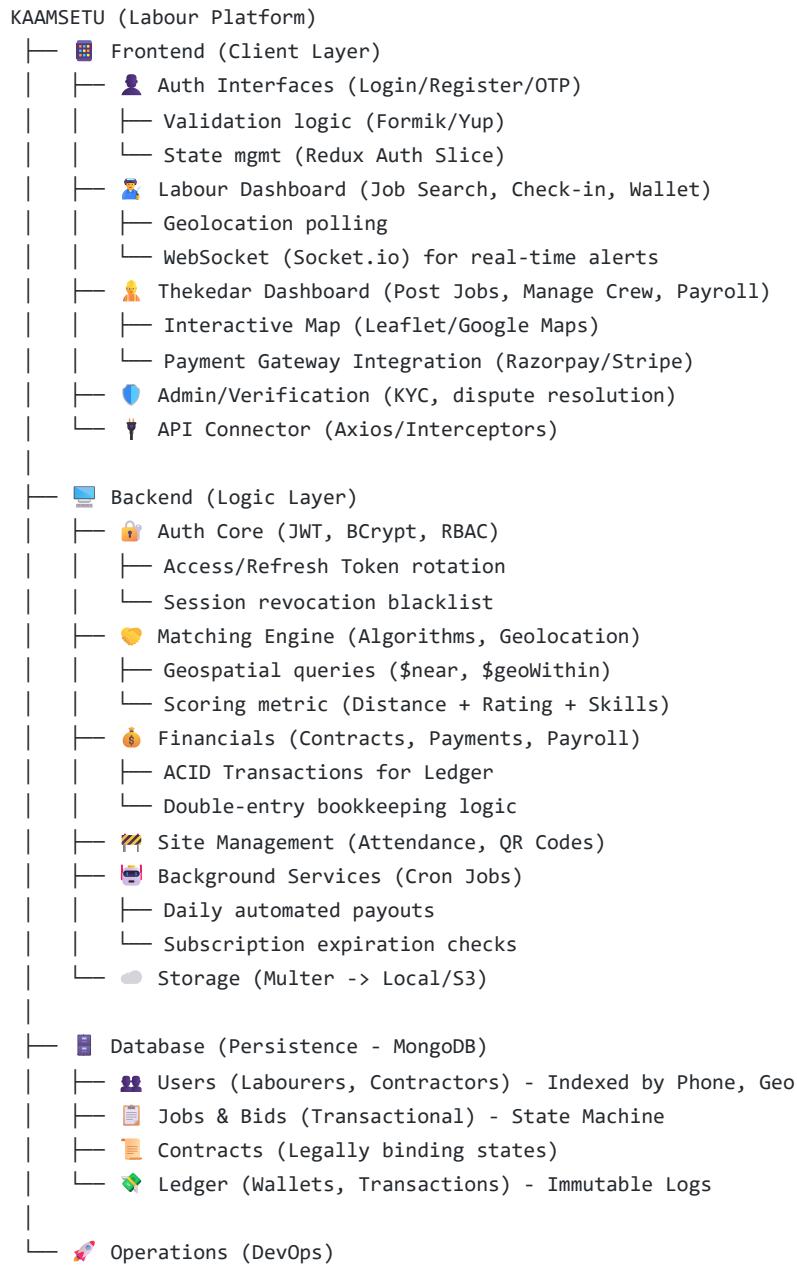


# KAAMSETU (MAJDOOR) - ARCHITECTURAL BLUEPRINT

**Version:** 5.1 (Interview Ready) **Role:** Principal Architect Documentation **Target Audience:** Engineering Team & Interview Panel

## 1 MIND GRAPH / KNOWLEDGE MAP

### "The Nervous System of KaamSetu"



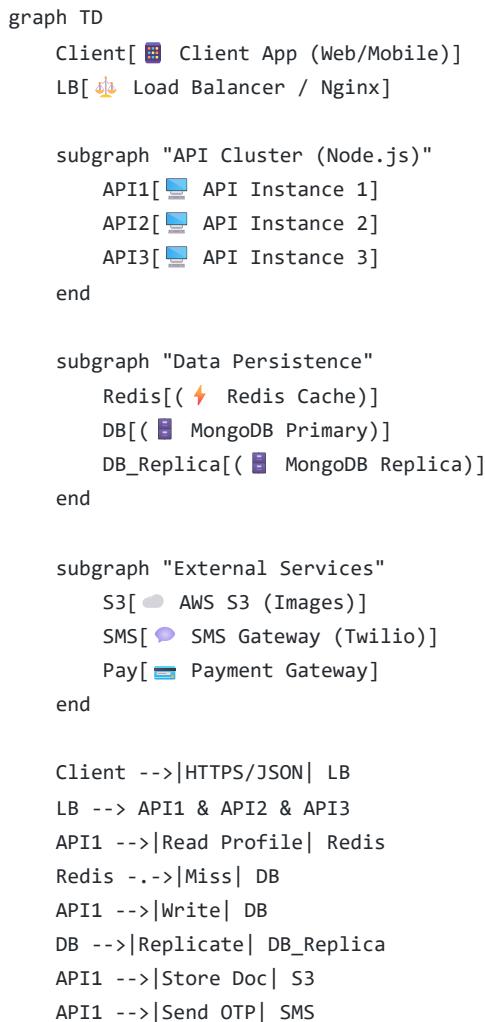
- └ CI/CD (GitHub Actions)
- └ Cloud (Render/AWS)
- └ Security (Helmet, RateLimit, XSS, HPP)
- └ Monitoring (Morgan, Winston, Sentry)

#### Node Explanation:

- **Frontend State:** Used Redux Toolkit to prevent prop-drilling in deep widget trees (e.g., Wallet Balance updating when a job is finished).
- **Matching Engine:** Decoupled from the main API. Can be spun out as a microservice if load increases.
- **Ledger:** Designed as an "Append Only" log system. We never overwrite a transaction balance; we only add new credit/debit entries.

## 2 SYSTEM & BLOCK DIAGRAMS

### A) High-Level Architecture (The 10,000ft View)



## B) Authentication Flow (JWT + Refresh Token)

```
sequenceDiagram
    participant User
    participant Client
    participant API
    participant DB

    User->>Client: Enter Credentials
    Client->>API: POST /auth/login
    API->>DB: Find User & Compare Hash
    DB-->>API: User Valid
    API->>API: Generate AccessToken (15m) & RefreshToken (7d)
    API-->>Client: Return AccessToken + Set HttpOnly Cookie (RefreshToken)

    note right of Client: Time Passes (16 mins later)

    Client->>API: GET /profile (w/ Expired AccessToken)
    API-->>Client: 401 Unauthorized
    Client->>API: POST /auth/refresh (Cookie sent automatically)
    API->>DB: Verify RefreshToken in DB Checksum
    API-->>Client: New AccessToken
```

## C) Feature Execution: "Thekedar Hires Labourer"

```
flowchart TD
    A[Thekedar clicks 'Hire'] --> B{Balance Sufficient?}
    B -- No --> C[Redirect to Payment Gateway]
    B -- Yes --> D[API: POST /contracts/create]
    D --> E[Backend: Start Transaction]
    E --> F[Lock Job Status to 'Filled']
    E --> G[Create Contract Record]
    E --> H[Escrow Amount from Wallet]
    E --> I[Commit Transaction]
    I --> J[Notify Labourer Via Socket]
    J --> K[Update UI]
```

## 3 AGENT-DECIDED TECH STACK

### "Rationale for Production Choices"

#### Frontend: React 19 + Redux Toolkit

- **Why?** React 19 minimizes re-renders with the new Compiler, essential for low-end devices often used by our target demographic (Labourers).
- **Why Redux Toolkit?** We handle complex synchronous global state: `User Profile` + `Wallet Balance` + `Active Bids` + `GPS Location`. Context API triggers too many re-renders. RTK is efficient.

- **Alternatives Considered:**

- *Next.js*: Good for SEO, but we are building a closed platform (Dashboard) where SEO matters less than app-like interactivity.
- *Vue*: Easier curve, but React ecosystem (ecosystem for Maps, Payments) is richer.

## **Backend: Node.js + Express 5**

- **Why Node?** Validating GeoJSON data is CPU light but IO heavy. Node's Event Loop handles thousands of concurrent location pings better than threaded Python/Ruby.
- **Why Express 5?** Native Promise support in routing. Removes the need for `asyncHandler` wrappers or `try-catch` hell in every controller.
- **Alternatives Considered:**
  - *Go (Golang)*: Better raw performance, but development speed is slower. We needed rapid iteration.
  - *NestJS*: Too much boilerplate for an MVP. Express is unopinionated and fast to set up.

## **Database: MongoDB (Mongoose)**

- **Why NoSQL?**

1. **Polymorphism:** A `User` can be a Labourer (Skills, Rating) or Thekedar (Company Name, GST). Single collection with flexible schema handles this best.
2. **Geospatial:** MongoDB's `$near` operator is industry standard for location queries.

- **Why Mongoose?**

- Validation: We need to ensure `walletBalance` never drops below 0 via schema validators (`min: 0`).
- Relationships: Virtual populate allows us to link `Bids` to `Jobs` without complex SQL joins.

## **Security & Infra**

- **Files:** Multer streams files to disk/cloud. We check magic numbers (file signatures) to prevent spoofed executables being uploaded as `.jpg`.
- **API Security:**
  - `Helmet` : Sets HTTP headers (Strict-Transport-Security, X-Frame-Options).
  - `HPP` : Prevents HTTP Parameter Pollution (e.g. `?id=1&id=2` attacks).
  - `RateLimit` : 100 req/15min per IP.

## **4 COMPLETE FILE & FOLDER STRUCTURE**

```
/labour-platform
├── /backend
│   ├── package.json      # Deps: express, mongoose, bcryptjs, jsonwebtoken
│   ├── server.js         # Entry point, Middleware registration, DB Connect
│   ├── /config
│   │   └── db.js          # Mongoose connection logic with retry strategy
│   ├── /controllers       # 🌸 Controller Logic (Req -> Business Logic -> Res)
│   │   ├── authController.js    # Login, Register, Refresh
│   │   ├── jobController.js    # CRUD Jobs, GeoSearch
│   │   └── paymentController.js # Wallet logic, Gateway callbacks
│   ├── /middleware         # 💡 Interceptors
│   │   ├── authMiddleware.js  # Decodes JWT, attaches req.user
│   │   └── errorMiddleware.js # Centralized Error Handling (JSON resp)
```

```

|   |   └── uploadMiddleware.js # Multer config (Size limit, File filter)
|   ├── /models           # 📁 Data Schemas
|   |   ├── User.js        # index: { location: '2dsphere' }
|   |   ├── Job.js
|   |   └── Transaction.js
|   ├── /routes           # 🚦 API Routes
|   |   └── v1/             # Versioned routes
|   └── /utils             # 🛠️ Helpers
|       ├── geocoder.js   # Google Maps API wrapper
|       └── sendEmail.js  # SMTP wrapper
|
└── /frontend
    ├── package.json
    ├── /public          # index.html, manifest.json (PWA)
    └── /src
        ├── /assets
        ├── /components     # 🏙️ Shared UI
        |   ├── /common      # Buttons, Inputs, Modals
        |   └── /layout       # Navbar, Sidebar, Footer
        ├── /features        # 🎨 Redux Slices & Thunks
        |   ├── auth/         # authSlice.js, authService.js
        |   └── jobs/         # jobSlice.js
        ├── /hooks           # 💡 Custom Hooks (useGeoLocation, useAuth)
        ├── /pages           # 📄 Route Views
        |   ├── Dashboard.js
        |   └── Login.js
        ├── /services         # 🕵️ API Connectors (Axios instances)
        └── App.js           # Router Config

```

## 5 DATABASE DESIGN & DATA FLOW

### ER Diagram (Textual Representation)

#### User Collection

```
{
  "_id": "ObjectId('...')",
  "name": "Raju Mistry",
  "phone": "+919876543210",
  "role": "labour",
  "password": "$2a$10$hash...",
  "location": {
    "type": "Point",
    "coordinates": [77.1025, 28.7041] // [Long, Lat]
  },
  "skills": ["Mason", "Tiles"],
  "walletBalance": 1500.00,
  "isVerified": true,
```

```
"createdAt": "ISODate(...)"  
}
```

#### Indexes:

1. phone : Unique. Fast lookups during login.
2. location : 2dsphere . Enabling \$near queries.

#### Job Collection

```
{  
  "_id": "ObjectId('...')",  
  "postedBy": "Ref(User)",  
  "title": "Need 5 Masons for 3 Days",  
  "wagePerDay": 800,  
  "requiredSkills": ["Mason"],  
  "location": { "type": "Point", "coordinates": [...] },  
  "status": "open", // open, filled, completed, cancelled  
  "applicants": ["Ref(User)", "Ref(User)"]  
}
```

#### Transaction Collection (The Ledger)

```
{  
  "_id": "ObjectId('...')",  
  "user": "Ref(User)",  
  "amount": 500,  
  "type": "credit", // credit, debit  
  "reason": "Job Payment - Contract #123",  
  "balanceAfter": 2000,  
  "metadata": { "contractId": "..." }  
}
```

## Scaling Logic

- **Sharding:** We shard the User collection based on zipCode or broad region . Labour markets are hyper-local; a worker in Delhi doesn't need to match with a job in Mumbai.
- **Archiving:** Jobs older than 6 months are moved to Jobs\_Archive collection to keep the hot index small.



6

## FEATURE-BY-FEATURE DEEP INTERNAL LOGIC

### Feature 1: The "Find Labour" Algorithm

**Objective:** Return the most relevant workers for a specific job post.

**API Endpoint:** POST /api/jobs/match

**Algorithm Steps:**

1. **Input Parsing:** Extract Job Lat/Lon and Required Skills.

2. **Geo-Query:**

```
User.find({
  role: 'labour',
  skills: { $in: jobSkills },
  location: {
    $near: {
      $geometry: { type: "Point", coordinates: [lon, lat] },
      $maxDistance: 5000 // 5km radius
    }
  }
})
```

3. **Ranking (In-Memory or Aggregation):**

- o Boost score if `isVerified` is true.
- o Boost score if `rating > 4.5`.
- o Penalize if `lastActive > 7 days ago`.

4. **Response:** Return top 20 User Objects (sanitized, no passwords).

## Feature 2: Atomic Wallet Transfer

**Objective:** Move money from Thekedar to Labourer safely.

**API Endpoint:** POST /api/payments/payout

**Logic (Mongoose Session):**

```
const session = await mongoose.startSession();
session.startTransaction();
try {
  const employer = await User.findById(employerId).session(session);
  if (employer.walletBalance < amount) throw new Error('Insufficient Funds');

  // Debit Employer
  employer.walletBalance -= amount;
  await employer.save();

  // Credit Worker
  const worker = await User.findById(workerId).session(session);
  worker.walletBalance += amount;
  await worker.save();

  // Create Logs
  await Transaction.create([{ type: 'debit', ... }], { session });
  await Transaction.create([{ type: 'credit', ... }], { session });

  await session.commitTransaction();
} catch (error) {
  await session.abortTransaction(); // ROLLBACK EVERYTHING
}
```

```
    throw error;  
}
```



7

## AUTH, SECURITY & COMPLIANCE

### Password Security

- **Algorithm:** BCrypt (Blowfish cipher).
- **Salt Rounds:** 10 (Standard compromise between security and speed).
- **Why not MD5/SHA?** They are fast, making them vulnerable to Rainbow Table attacks. BCrypt is slow by design.

### JWT Strategy

- **Access Token:** Holds `userId` and `role`. Signed with `JWT_SECRET`. Expiry: 15 mins.
- **Refresh Token:** Random String stored in DB with `userId`. Expiry: 7 days.
- **Flow:** When Access Token expires (401), frontend hits `/refresh`. Backend checks if Refresh Token in cookie matches DB. If yes, issue new Access Token.
- **Revocation:** If user clicks "Logout all devices", we delete the Refresh Token from DB. All sessions die instantly.

### Common Attack Prevention

- **NoSQL Injection:** `req.body.username = { "$ne": null }` allows login without password.
  - *Fix:* `express-mongo-sanitize` strips keys starting with `$`.
- **XSS (Cross Site Scripting):** Malicious scripts in Job Description.
  - *Fix:* `xss-clean` middleware sanitizes HTML input. React escapes content by default.



8

## DEPLOYMENT, CI/CD & ENV SETUP

### Dockerfile Strategy

```
# Stage 1: Build Frontend  
FROM node:18 as build-stage  
WORKDIR /app/frontend  
COPY frontend/package*.json ./  
RUN npm install  
COPY frontend/ .  
RUN npm run build  
  
# Stage 2: Serve Backend  
FROM node:18  
WORKDIR /app  
COPY backend/package*.json ./  
RUN npm install --production  
COPY backend/ .  
# Copy frontend build to backend static folder
```

```
COPY --from=build-stage /app/frontend/build ./public
EXPOSE 5000
CMD ["node", "server.js"]
```

## CI Pipeline (GitHub Actions)

1. **Triggers:** Push to `main` or Pull Request.
2. **Lint:** Run ESLint.
3. **Test:** Run `npm test` (Jest unit tests).
4. **Deploy:** Only if tests pass, push Docker image to Container Registry (Docker Hub / ECR).
5. **Release:** Trigger Render/AWS to pull new image.



9

## DESIGN DECISIONS & TRADE-OFFS

### 1. Monolith vs Microservices

**Choice:** Monolithic Architecture (Modular).

- **Why:** Team size is small (< 5 engineers). Microservices introduce "Network Latency", "Distributed Tracing", and "Deployment Complexity" overhead.
- **Trade-off:** If the "Matching Service" crashes due to OOM, the "Auth Service" also goes down. We mitigate this with PM2 clustering.
- **Breaking Point:** When we reach >50 engineers or distinct scaling needs (e.g., Video Streaming feature needs distinct hardware).

### 2. Polling vs WebSockets

**Choice:** WebSockets (Socket.io) for Job Alerts.

- **Why:** Real-time is critical. A labourer needs to know about a job *instantly* to bid. Polling every 5 seconds wastes bandwidth and battery.
- **Trade-off:** Maintains open TCP connections, costing server RAM.
- **Mitigation:** Fallback to Long-Polling if connection unstable.

### 3. Consistency vs Availability (CAP Theorem)

**Choice:** CP (Consistency & Partition Tolerance) for financial data. AP (Availability) for Job Search.

- **Why:** Wallet balance MUST be consistent (CP). Job search results can be slightly stale (AP) without catastrophic failure.



10

## FLASH NOTES & REVISION CARDS



### FLASH CARD: Scalability

**Q: How do you handle 1 Million Users?** **A:** Horizontal Scaling (Auto-Scaling Groups), DB Sharding (Geo-based), Redis Caching for Read-Heavy data, and CDN for static assets.



### FLASH CARD: Docker

**Q: Difference between Image and Container?** **A:** Image is the blueprint (Class). Container is the running instance (Object).

### FLASH CARD: REST vs GraphQL

**Q: Why REST for this project?** **A:** Simple caching (browser/CDN), standard error codes (404, 500), and we don't have deeply nested data requirements that necessitate GraphQL's query flexibility.

### FLASH CARD: Indexing

**Q: Explain Compound Index.** **A:** An index on multiple fields, e.g., `{ location: "2dsphere", role: 1 }`. Crucial for queries like "Find Labourers near Coordinates". Order matters (Equality -> Sort -> Range).

---

## 1 1 COMPLETE TECHNICAL INTERVIEW Q&A

This section is designed to be **memorized** for the technical round.

### BEGINNER (Foundational Knowledge)

**1. Q: What is the difference between `dependencies` and `devDependencies` in `package.json`?**

- **A:** `dependencies` are required for the app to run in production (e.g., React, Express). `devDependencies` are only needed for development/building (e.g., Jest, ESLint, Nodemon).

**2. Q: Explain the React Virtual DOM.**

- **A:** It's a lightweight copy of the actual DOM. When state changes, React updates the Virtual DOM, diffs it with the previous version, and only updates the *changed* elements in the real DOM (Reconciliation), boosting performance.

**3. Q: usage of `useEffect` ?**

- **A:** It handles side effects in functional components, like data fetching, subscriptions, or manual DOM manipulation. It runs after render. Dependencies array controls when it re-runs.

**4. Q: Differences between `let` , `const` , and `var` ?**

- **A:** `var` is function-scoped and hoisted. `let` and `const` are block-scoped. `const` cannot be reassigned (though objects are mutable), while `let` can be.

**5. Q: What is Middleware in Express?**

- **A:** Functions that execute during the request-response cycle. They have access to `req` , `res` , and `next` . Used for logging, auth, parsing bodies, and error handling.

**6. Q: SQL vs NoSQL?**

- **A:** SQL (Relational) is table-based, has fixed schema, good for complex joins/transactions. NoSQL (Document/Key-Value) is flexible, scalable, and good for unstructured data.

**7. Q: What is a Promise?**

- **A:** An object representing the eventual completion (or failure) of an asynchronous operation. States: Pending, Fulfilled, Rejected.

**8. Q: Why use `async/await` compared to `.then()` ?**

- **A:** It makes asynchronous code look and behave like synchronous code, improving readability and making error handling (try/catch) cleaner.

#### 9. Q: What is Prop Drilling?

- **A:** Passing data from a parent to a deeply nested child through intermediate components that don't need the data. Solved via Context API or Redux.

#### 10. Q: HTTP Verbs meaning?

- **A:** GET (Retrieve), POST (Create), PUT (Replace), PATCH (Update partial), DELETE (Remove).
- 

## 🟡 INTERMEDIATE (Application Logic & Design)

#### 11. Q: How do you manage Global State in this application?

- **A:** I used **Redux Toolkit**. It centralizes state (User, Wallet, Jobs). I use `createSlice` for reducers and `createAsyncThunk` for async API calls, ensuring a predictable state container.

#### 12. Q: How does the Authentication System work here?

- **A:** It uses **JWT**. On login, server sends an Access Token (JSON body) and a Refresh Token (HttpOnly Cookie). Access token authorizes requests via Middleware. When it expires, the frontend silently calls `/refresh` to get a new one using the cookie.

#### 13. Q: Explain the Matching Algorithm you implemented.

- **A:** It's a Geospatial query using MongoDB's `$near` operator. It takes the Job's coordinates and searches the `User` collection for workers within 5km radius who have the matching `skill` in their skills array.

#### 14. Q: How do you handle File Uploads?

- **A:** Using **Multer** middleware. It handles `multipart/form-data` requests. I configured it to filter for image/pdf MIME types only and limited file size to 2MB to prevent server overload.

#### 15. Q: What is the purpose of `useEffect` cleanups?

- **A:** To prevent memory leaks. e.g., when a component unmounts, we must clear intervals, remove event listeners, or close WebSocket connections initiated in that effect.

#### 16. Q: How do you prevent double-booking of a job?

- **A:** Database Locking or Atomic Updates. Optimistic approach: `Job.findOneAndUpdate({ _id: jobId, status: 'open' }, { status: 'filled', contractor: id })`. If it returns null, someone else took it.

#### 17. Q: Explain CORS and how you handled it.

- **A:** Cross-Origin Resource Sharing. Browsers block requests from Domain A (Frontend) to Domain B (Backend). I used the `cors` middleware in Express to whitelist my Frontend's origin (e.g., localhost:3000 or mydomain.com).

#### 18. Q: What is an Environment Variable and why use it?

- **A:** Configuration outside code (e.g., API Keys, DB URI). It keeps secrets out of Git and allows different configs for Dev, Test, and Prod. Managed via `.env` and `dotenv`.

#### 19. Q: How do you debug a slow API endpoint?

- **A:** 1. Check logs (Morgan). 2. Use `console.time` / APM tools to spot bottlenecks. 3. Check DB Query plan (Need indexing?). 4. Check for blocking synchronous code.

## 20. Q: Difference between Authentication vs Authorization?

- **A:** Authentication = "Who are you?" (Login). Authorization = "What are you allowed to do?" (Permissions/Roles).

---

## ● ADVANCED (Scaling, Architecture, Security)

### 21. Q: Design a system to handle 10k concurrent job postings.

- **A:** 1. **Load Balancer** (Nginx) to distribute traffic. 2. **Queue** (RabbitMQ/Bull) to process job creation asynchronously if it involves heavy notifications. 3. **DB Sharding** to distribute writing load. 4. **Caching** is less useful for *writes*, but essential for *reads*.

### 22. Q: How would you secure the payment API?

- **A:** 1. **Idempotency Keys** to prevent double charges on retry. 2. **Signature Verification** (HMAC) for webhooks from Gateway to ensure authenticity. 3. **HTTPS** enforcement. 4. **RBAC** to ensure only the wallet owner initiates transfer.

### 23. Q: Explain Database Indexing strategies used.

- **A:** Used `2dsphere` for Location. Single Field Index on `phone` (Login). Compound Index on `status` + `createdAt` for filtering Job feeds efficiency. Indices trade write speed/storage for read speed.

### 24. Q: How does Node.js handle high concurrency if it's single-threaded?

- **A:** It uses the **Event Loop** and **Libuv**. The main thread is non-blocking. I/O operations (DB, Network) are offloaded to system kernel threads. When done, a callback is pushed to the Queue, which the Event Loop executes. Perfect for I/O heavy apps like ours.

### 25. Q: What is a Race Condition? How did you fix it in the Wallet?

- **A:** Two processes modifying data simultaneously causing inconsistency (e.g., Check Balance -> (Pass) -> Deduct). A parallel request could deduct again before the first saves. Fixed using **MongoDB Transactions** (ACID) or `findAndModify` atomic operations.

### 26. Q: Microservices vs Monolith - Why did you choose Monolith?

- **A:** Complexity trade-off. Microservices add distributed system problems (Network failure, eventual consistency). For an MVP/Mid-scale, a **Modular Monolith** (Clean separation of folders) gives 80% of the benefits (organization) with 0% of the network overhead.

### 27. Q: Critical Rendering Path optimization?

- **A:** 1. Code Splitting (React.lazy) to load bundles only when needed. 2. Image Optimization (WebP). 3. Minification of JS/CSS. 4. Preloading critical fonts.

### 28. Q: How do you handle DDoS attacks?

- **A:** Infrastructure layer: Cloudflare/AWS Shield. App layer: `express-rate-limit`. Server layer: Nginx connection limits.

### 29. Q: Explain Blue-Green Deployment.

- **A:** Two identical environments (Blue=Live, Green=Idle). Deploy new version to Green. Run tests. Switch Router to point to Green. If error, switch back to Blue instantly (Rollback).

### 30. Q: What happens when you type a URL and hit Enter? (Deep Dive)

- **A:** 1. DNS Resolution (Domain -> IP). 2. TCP Handshake. 3. TLS Handshake (HTTPS key exchange). 4. HTTP Request sent. 5. Server processes (LB -> Node -> DB). 6. HTTP Response (HTML). 7. Browser parses DOM, fetches assets, renders pixels.



## HOW TO EXPLAIN THIS PROJECT

**2-Minute Pitch (The Hero's Journey):** "I built KaamSetu to solve a specific problem in the unorganized labor market: mismatch and trust. Traditionally, labourers stand at checks waiting to be picked up. My platform digitizes this.

I architected it using a MERN stack. The core is the **Matching Engine** which uses Geospatial queries to link workers and contractors instantly. I prioritized performance using React 19 and Node js, and security using JWT and Helmet.

The most challenging part was the **Financial Ledger**. Ensuring that wallet transactions were atomic and error-proof required implementing MongoDB ACID transactions.

Today, the system is fully containerized with Docker and features a CI/CD pipeline, ready for real-world deployment."



## FUTURE IMPROVEMENTS

1. **Server-Side Rendering (SSR) with Next.js:** To improve initial load time and SEO for public job public link sharing.
2. **Redis Caching Layer:** To cache 'Hot Jobs' and user profiles, reducing DB hits by 80%.
3. **BullMQ:** To handle notification broadcasting (SMS/Email) in the background so the API remains snappy.

*Created by Antigravity Autonomous Agent | 2026*